

Experience	NVIDIA · Performance Software Engineering Intern Aug 2020 – Present <ul style="list-style-type: none">Optimizing sparse BERT inference performance for TensorRT in C++, enabling a potential 50% reduction in inference time, memory usage, and power usage for customers
	Uber ATG · Research Intern Jan 2020 – Aug 2020 <ul style="list-style-type: none">Improved object detection by 90% (AP) and motion forecasting by 22% (L2) of a self-driving neural net under realistic positional error, significantly improving safety for future ridersWrote a first author paper on the learned positional error correction system (accepted at CoRL)
	Google Brain · Software Engineering Intern May 2019 – Aug 2019 <ul style="list-style-type: none">Unlocked K-FAC for over 370,000 users by implementing and open sourcing automatic support for arbitrary neural network architectures and integrating it into the Keras ecosystem 🔗Enabled simple multi-node, multi-GPU/TPU training for users by incorporating TensorFlow's Distribution Strategy and efficient distributed operation placementDesigned, created, and open-sourced idiomatic, reproducible training recipes for users while carefully considering hyperparameter ranges, baselines, datasets, and models 🔗
	John Hancock Financial · Data Science Intern May 2018 – Aug 2018 <ul style="list-style-type: none">Achieved a fraud detection rate of 63% through designing an unsupervised ML modelDeployed 25 fraud identifying rules in SQL that correctly flagged 100+ out of 20,000+ claims
	Sunnybrook Research Institute · Software Developer Intern Jul 2017 – Aug 2017 <ul style="list-style-type: none">Improved MRI segmentation accuracy by up to 80% and reduced time to contour MRI scans from ~5 hrs to ~40 mins by implementing techniques including watershed and clustering
Publications	Nicholas Vadivelu , Mengye Ren, James Tu, Jingkang Wang, Raquel Urtasun. Learning to Communicate and Correct Pose Errors. In <i>Conference on Robotics Learning (CoRL)</i> , Virtual, 2020 Pranav Subramani, Nicholas Vadivelu , Gautam Kamath. Enabling Fast Differentially Private SGD via Just-in-Time Compilation and Vectorization. In <i>NeuRIPS Privacy-Preserving Machine Learning Workshop</i> , Virtual, 2020 🔗
Open Source	PyTorch Ignite : Improved performance by up to 63% by designing and implementing async updates for distributed metrics with tests and documentation 🔗
Projects	Thrive Life Simulator : Created a 3D ray-casting game engine from scratch for a dinosaur world simulation game in Java with object-oriented design and detailed documentation 🔗 Vim Clone : Recreated the text editor using object-oriented design and C++ best practices, such as implementing the Model-View-Controller pattern and extensively using STL functionality
Leadership	Data Science Club Lectures : Designed and presented workshops about neural networks in TensorFlow , machine learning in scikit-learn , and data cleaning in pandas for 300+ students 🔗 WATonomous Design Team : Implemented real-time object detection in Tensorflow , OpenCV
Education	University of Waterloo · Computer Science & Statistics (B. Math) Cumulative GPA: 3.94/4.00 - Dean's List 2017 – 2022